

Round 2

AI/ML Take-Home Task - 48h

Task Overview:

Goal: show clear thinking, sensible trade-offs and clean, testable endpoints.
(Note: Heavy modeling is not required.)

Using the provided sample data, build a mini-pipeline that:

Pre-Transcription (*implementation not required - outline only*):

Briefly outline what audio pre-processing (if any) you would consider and why.

1. **Transcription:**

Use Whisper or similar. Keep timestamps if convenient.

2. **Topic Extraction:**

Extract the top 5 discussion topics from the transcript.

Optionally, you may also consider other conversational/social cues (*explain your approach, only topics must be implemented*).

3. **Vectorization:**

Choose approach for extracted topics (one-hot, TF-IDF, or embeddings).

4. **Load User Profiles:**

Use provided JSON psychometric profiles. Each user represents a speaker in the audio file: User 1 = Joe Rogan; User 2 = Elon Musk.

5. **Fusion: topics x personality (open-ended)**

Propose and justify how you weight topics relative to each user's personality before combining. Construct a combined vector per user.

6. **Compatibility Score + Interpretation:**

Compute a compatibility score between the two users (e.g. cosine similarity over their combined vectors) and returns an interpretation of that score.

7. **Fast API Endpoints:**

Wrap your pipeline into a FastAPI application that exposes:

POST/transcribe (returns a transcript), POST/summarise (returns topic list) and POST/match (returns compatibility score).

8. **Live, Testable Demo:**

Make the app reachable so it can be tested via browser (Swagger/OpenAPI UI).

Sample Data (in sample_data/):

Audio: 'Is Mars the Future of Humanity?' - Joe Rogan Asks Elon Musk

JSON:

```
Traits = [openness, conscientiousness, extraversion, agreeableness, neuroticism]
[
  { "id": "user_1", "psychometrics": [0.8, 0.4, 0.7, 0.2, 0.9] },
  { "id": "user_2", "psychometrics": [0.3, 0.9, 0.1, 0.6, 0.4] }
]
```

Suggested Tech Stack:

Use any combination of:

Transcription: OpenAI Whisper, Hugging Face whisper, AssemblyAI

Topic Extraction: spaCy, Gensim, Transformers (BERT/DistilBERT), keyBERT

Vectorization: scikit-learn TF-IDF, SentenceTransformers, et

Matching Logic: cosine similarity on combined vectors (numpy / scikit-learn)

LLM Fine Tuning: GPT-2, LLaMA via Hugging Face transformers + accelerate or OpenAI API

Orchestration: PyTorch/TensorFlow; LangChain (optional)

Deliverables:

1. GitHub repo with clean, dated, well-documented code.
2. README covering:
 - How to load and use the sample data
 - Architecture & design decisions
 - Pre-Processing Outline (if any)
 - Topic vectorization method and any considerations of social cues.
 - Fusion/Weighting Rationale
 - Matching logic and edge-case handling (thresholds or re-sampling)
 - Live URL
3. Short write-up (≤ 300 words) on: reasoning, trade-offs and next steps.